


Apache Flink v2.1 Kickstarter

This project demonstrates the **real-world power of Apache Flink**—a distributed stream processing engine built for **low-latency**, **high-throughput**, **and exactly-once stateful computation**. As a core component of the **signalRoom technology stack**, it highlights how Flink's advanced features—such as event-time semantics, checkpointing, watermarks, and dynamic scaling—can be leveraged to build resilient, production-grade streaming pipelines. Staying true to our mission of **empowering developers through practical**, **open knowledge**, this project extends the concepts from Confluent's "Building Apache Flink Applications in Java" series into runnable, end-to-end examples that showcase how to design, deploy, and optimize Flink jobs for real business impact.

Embarking on this journey, you'll explore how Flink seamlessly integrates with modern data infrastructures, including **Confluent Cloud for Apache Flink**, **AWS Secrets Manager**, **AWS Systems Manager Parameter Store**, and **Apache Iceberg tables in Snowflake**. This project not only illustrates Flink's capabilities in handling complex event processing and state management but also demonstrates best practices for secure configuration management and scalable data storage.



Flink App

Description

Flink App

Description

Data Generator App This Apache Flink application simulates realistic flight activity for the fictional airlines
Sunset Air and Sky One Airlines, showcasing seamless integration across modern
data streaming technologies. Using Flink's high-throughput processing capabilities,
flight events are continuously published to dedicated Kafka topics
(airline.sunset and airline.skyone), enabling real-time monitoring,
transformation, and analytics. In parallel, the same synthetic data is written to
Apache Iceberg tables (apache_kickstarter.airlines.sunset_airline
and apache_kickstarter.airlines.skyone_airline) stored in AWS S3,
creating a unified, scalable foundation for both streaming and historical analysis. Built
entirely in Java, this implementation overcomes current PyFlink limitations—namely,
the absence of a Python-native data generator source—while demonstrating best
practices for stateful stream processing and unified batch/stream architecture. The
result is a powerful, extensible data pipeline that bridges real-time event
processing and analytical persistence, illustrating how Flink can serve as the
backbone of next-generation streaming systems.

Flight Consolidator App

This Apache Flink application ingests flight event data from the airline.sunset and airline.skyone Kafka topics, transforms and standardizes the records into a unified airline.flight Kafka topic, and persists them to the apache_kickstarter.airlines.flight Apache Iceberg table. Implemented in both Java and Python, it demonstrates cross-language stream processing, schema unification, and seamless integration between Kafka, Flink, and Iceberg for end-to-end real-time and analytical data workflows.

Flyer Stats App

This Apache Flink application consumes flight event data from the airline.flight Kafka topic, performs real-time aggregations to compute flyer-level statistics, and publishes the results to the airline.flyer_stats Kafka topic as well as the apache_kickstarter.airlines.flyer_stats Apache Iceberg table. Implemented in both Java and Python, it demonstrates how to build unified, stateful streaming and analytical pipelines for continuous aggregation and historical insight.

Originally created by Wade Waldron, Staff Software Practice Lead at Confluent Inc., these applications are built to showcase Apache Flink's enterprise-grade capabilities through realistic, production-scale scenarios. The workflow begins by securely retrieving Kafka Cluster and Schema Registry API keys from AWS Secrets Manager, followed by loading essential Kafka client configuration parameters from AWS Systems Manager Parameter Store. This design highlights best practices for secure, automated configuration management and real-world Flink integration across AWS and Confluent Cloud environments.

We go beyond simple transformation and enrichment — this project **streams data into Kafka topics and** *Apache Iceberg Tables*, purpose-built for **high-performance**, **large-scale analytics**. Iceberg brings **ACID compliance**, schema evolution, and **highly scalable**, **real-time data processing with durable storage**, forming the analytical backbone of this architecture.

Our exploration of **Apache Flink** extends beyond Java into **Python-based applications**, showcasing the power of **Flink SQL**, **Table API**, and **DataStream API**. Each serves a distinct purpose:

- Flink SQL for declarative, real-time analytics
- Table API for a blend of SQL expressiveness and programmatic control
- DataStream API for efficient, event-driven stream processing

Together, they deliver a **unified**, **flexible**, **and scalable approach** to solving complex streaming challenges —bridging real-time and batch processing with elegance and performance.

Run the Docker containers (fully compatible with both **Apple Silicon's M-series** and **x86 architectures**) and experience firsthand how Apache Flink can power **robust**, **enterprise-grade streaming applications** that are as practical as they are powerful.

Table of Contents

- 1.0 The apache_flink-kickstarter Just Got Smarter: From Stream Processing to Real-Time Intelligence
 - 1.1 Apache Flink v2.1: Real-Time Al and Smarter Streaming
 - 1.2 Apache Iceberg v1.10: A Lakehouse Built for Modern Workloads
 - 1.3 Behind the Scenes: Code, Performance, and Best Practices
 - 1.4 Why It Matters
- 2.0 Are these examples better described as Flink Jobs or Flink Applications—and why does the distinction matter?
- 3.0 Let's get started!
 - o 3.1 What You'll Learn
 - 3.2 Beyond Kafka: Iceberg Integration
 - o 3.3 Run It Locally Fast
 - o 3.4 Dive In
 - o 3.5 Getting Started
 - 3.5.1 Cloud Prerequisites
 - 3.5.2 Local Prerequisites
 - 3.5.3 Clone the Repository
 - 3.5.4 Set Up Your Environment
 - 3.5.5 Deploy Flink
 - 3.5.6 You're Ready!
 - 3.6 DevOps in Action: Running Terraform
 - 3.6.1 Deploy Terraform
 - o 3.7 DevOps in Action: Running Terraform in the cloud
 - 3.7.1 Run from the cloud
 - 3.8 Visualizing the Terraform Configuration
- 4.0 Hands-On Kickoff: Practical Examples for Rapid Learning
 - 4.1 Flink Applications Powered by Java on a locally running Apache Flink Cluster in Docker
 - 4.2 Flink Applications Powered by Python on a locally running Apache Flink Cluster in Docker
 - 4.3 Flink Applications Powered by Python on Confluent Cloud for Apache Flink
- 5.0 Powering Analytics with Apache Iceberg tables in Snowflake
- 6.0 Resources

1.0 The apache_flink-kickstarter Just Got Smarter: From Stream Processing to Real-Time Intelligence

The apache_flink-kickstarter project has been completely refreshed to keep pace with the fast-moving evolution of **Apache Flink v2.1** and **Apache Iceberg v1.10**. This isn't just a routine maintenance update— it's a strategic move that advances your real-time data platform into the next era of streaming and Al-driven decision-making.

The upgrade updates the project with the latest APIs, performance improvements, and intelligent runtime hooks introduced since Apache Flink v1.20 and Apache Iceberg v1.7. The result: faster pipelines, improved observability, and new features that merge data engineering and data science.

1.1 Apache Flink v2.1: Real-Time Al and Smarter Streaming

If you've been holding steady on v1.20, the v2.1 release is the *inflection point* where **stream processing** meets machine learning:

- **Al/ML inside the stream:** Native model DDLs and inference functions allow you to deploy and run trained models directly in your Flink jobs—no sidecars, no hacks.
- Richer SQL & Table API: Handle complex and semi-structured data with new Process Table Functions (PTFs), faster joins, and more expressive queries.
- **Operational superpowers:** Enhanced runtime hooks enable detailed performance tuning, state inspection, and adaptive write strategies.
- **Evolution of purpose:** Flink's story is shifting from *real-time data processing* to *real-time intelligence*.

1.2 Apache Iceberg v1.10: A Lakehouse Built for Modern Workloads

For those on Apache Iceberg v1.7.x, version v1.10 marks the point when the open-table format truly matures:

- Broader ecosystem support: Native compatibility with the latest Flink and Spark versions for a seamless hybrid stack.
- Smarter table specification: Enables support for geospatial types, variant columns, and deletion vectors—enhancing advanced analytics and governance capabilities.
- **Operational finesse:** Smarter compaction, more precise stats, and streaming-friendly enhancements that simplify managing long-running jobs.
- Future-proofing your platform: Keeping it up-to-date guarantees access to the newest spec features and ongoing support from engine developers and vendors.

1.3 Behind the Scenes: Code, Performance, and Best Practices

The refreshed apache_flink-kickstarter codebase not only adopts new APIs but also incorporates modern engineering principles:

- Deprecated calls replaced with new unified APIs.
- Best practices for state management and checkpointing to ensure reliability.
- Optimized serialization, descrialization, and connector performance for Kafka, Iceberg, and S3.
- Cleaner configuration and modularized build logic for easier deployment across environments.

- Added thread-safety for random number generation in Flink applications.
- Improved overall **performance** in all Flink applications.

With these updates, the project isn't just a set of examples—it's a launchpad and a *blueprint*. It lets you spin up Flink + Iceberg pipelines in minutes and learn exactly how to architect streaming systems built for the scale and intelligence of 2025.

1.4 Why It Matters

The shift from batch analytics to continuous intelligence is occurring now. Apache Flink v2.1 and Apache lceberg v1.10 are not just incremental upgrades—they symbolize the merging of **streaming**, **AI**, and **lakehouse architecture**.

By upgrading today, you're not just boosting performance—you're securing your data platform for the next era of real-time, model-driven decision systems.

2.0 Are these examples better described as Flink Jobs or Flink Applications—and why does the distinction matter?

"What's in a name? That which we call a rose by any other name would smell just as sweet."

-- William Shakespeare

Did you know? In the Flink ecosystem, what many call "jobs" are actually better described as **Flink applications**—and that distinction matters.

A Flink application isn't just a single computation or stream transformation—it's a **self-contained**, **event-driven system** that coordinates state management, fault tolerance, and complex dataflows across distributed clusters. In essence, each one is a **mini software ecosystem**, designed to solve real-world data challenges at scale.

By calling them *applications*, we highlight their **depth, sophistication, and architectural completeness**. They don't just process data—they **model business logic, orchestrate stateful workflows, and power streaming intelligence** much like any robust production-grade software system.

(Curious why the terminology matters? Check out the rationale here!)

3.0 Let's get started!

As of October 4, 2024, Confluent Cloud for Apache Flink (CCAF) unlocks incredible potential for real-time stream processing—along with a few key nuances developers should know. Currently, Confluent Cloud for Apache Flink (CCAF) does not support the DataStream API, and its Table API has limited functionality (see supported features here), supporting Java and Python-based Flink applications. But don't worry—this repository bridges those gaps with practical, production-ready examples.

3.1 What You'll Learn

In this repo, you'll find **Java-based examples using the Flink DataStream API** to go beyond Flink's standard capabilities—like:

• Securely retrieving Kafka Cluster API keys from AWS Secrets Manager

• Dynamically loading **Kafka client configuration** from **AWS Systems Manager Parameter Store**—all directly integrated within a running Flink job.

Prefer Python? You're covered there too. Using the **Flink Table API** and **User-Defined Table Functions (UDTFs)**, you can achieve the same secure, dynamic configuration loading while staying fully in Python.

3.2 Beyond Kafka: Iceberg Integration

We don't just stream data into Kafka topics—this project also demonstrates how to **sink data into Apache Iceberg tables**, combining **real-time event streaming** with **durable**, **analytical storage** for hybrid batch/stream processing. By integrating Iceberg, you gain access to features like **ACID compliance**, **schema evolution**, and **time travel queries**—all essential for building robust, enterprise-grade data platforms that support both real-time and historical analytics.

3.3 Run It Locally — Fast

Spin up an **Apache Flink v2.1 cluster locally** in minutes using Docker—fully compatible with both **Apple Silicon's M-series** and **x86 machines**. The Flink cluster connects seamlessly to **Confluent Cloud** for Kafka and Schema Registry resources, giving you a realistic, end-to-end development environment.

3.4 Dive In

Explore the examples, run the containers, and see how to build **secure**, **scalable**, **enterprise-grade Flink applications** that push the limits of what's possible with modern data streaming. The future of real-time data pipelines starts here.

3.5 Getting Started

Follow these steps to set up your environment and launch the complete **Apache Flink Kickstarter** stack—locally or in the cloud.

3.5.1 Cloud Prerequisites

Before you begin, ensure you have access to the following cloud accounts:

- AWS Account with SSO configured
- Confluent Cloud Account for Kafka and Schema Registry resources
- **Docker Account** to pull and run containers
- GitHub Account with OIDC configured for AWS authentication
- Snowflake Account for data warehousing and analytics
- Terraform Cloud Account for automated infrastructure provisioning

3.5.2 Local Prerequisites

Make sure the following tools are installed on your local machine:

- AWS CLI version 2
- Confluent CLI version 4 or higher
- Docker Desktop

- Java JDK 21
- Python 3.12
- Terraform CLI version 1.13.4

3.5.3 Clone the Repository

```
git clone https://github.com/j3-signalroom/apache_flink-kickstarter.git
cd apache_flink-kickstarter
```

3.5.4 Set Up Your Environment

You can set up the complete environment in one of two ways:

- Option A: Configure and run Terraform Cloud locally.
- Option B: Use GitHub Actions to automatically create all resources.

Both methods provision a fully functional end-to-end data streaming architecture, including:

- A Confluent Cloud environment with a Kafka cluster and pre-configured example topics
- AWS Secrets Manager for secure Kafka API key storage
- AWS Systems Manager Parameter Store for Kafka client configuration (consumer/producer properties)
- An AWS S3 bucket with a dedicated warehouse/ folder serving as the landing zone for Apache lceberg Tables—populated by two Python-based Flink applications

3.5.5 Run Flink Locally

You can run **Apache Flink** locally on your Mac or use the **provided Docker containers** to launch Flink and lceberg seamlessly—compatible with both **Apple Silicon (M-series)** and **x86 architectures**.

3.5.6 You're Ready!

Once deployed, you'll have a **complete enterprise-grade data streaming environment**—secure, scalable, and ready to power both **real-time** and **analytical** use cases.

3.6 DevOps in Action: Deploying Terraform

Install the Terraform CLI on your local machine, and make sure you have an HCP Terraform account to run the Terraform configuration. Learn how to set up Terraform Cloud for local use by clicking here.

3.6.1 Deploy Terraform

<snowflake_warehouse></snowflake_warehouse>	snowflake-warehouse=
<service_account_user></service_account_user>	service-account-user=
<pre><3ERVICE_ACCOUNT_U3ER></pre>	day-count= <day_count></day_count>
latest>	auto-offset-reset= <earliest td="" <=""></earliest>
<number_of_api_keys_to_retain></number_of_api_keys_to_retain>	number-of-api-keys-to-retain=
<pre>admin-service-user-secrets-root- path=<admin_service_user_secrets_root_path></admin_service_user_secrets_root_path></pre>	

Argument placeholder	Replace with
<sso_profile_name></sso_profile_name>	your AWS SSO profile name for your AWS infrastructue that host your AWS Secrets Manager.
<confluent_api_key></confluent_api_key>	your organization's Confluent Cloud API Key (also referred as Cloud API ID).
<confluent_api_secret></confluent_api_secret>	your organization's Confluent Cloud API Secret.
<sn0wflake_warehouse></sn0wflake_warehouse>	the Snowflake warehouse (or "virtual warehouse") you choose to run the resources in Snowflake.
<service_account_user></service_account_user>	the Snowflake service account user who is to be assigned the RSA key pairs for its authentication.
<day_count></day_count>	how many day(s) should the API Key be rotated for.
<auto_offset_reset></auto_offset_reset>	Use earliest, when you want to read the first event in a Kafka topic. Otherwise, specify latest.
<number_of_api_keys_to_retain></number_of_api_keys_to_retain>	Specifies the number of API keys to create and retain.
<admin_service_user_secrets_root_path></admin_service_user_secrets_root_path>	the root path in AWS Secrets Manager for admin service user secrets.

To learn more about this script, click here.

3.7 DevOps in Action: Deploying Terraform in the cloud

In order to run the Terraform configuration from GitHub, the Terraform Cloud API token and Confluent Cloud API Key are required as GitHub Secret variables. Learn how to do to get the Terraform Cloud API token and Confluent Cloud API key here.

3.7.1 Run from the cloud


Follow these steps:

a. Deploy the Repository: Ensure that you have cloned or forked the repository to your GitHub account.

- b. **Set Required Secrets and Variables**: Before running any of the GitHub workflows provided in the repository, you must define at least the AWS_DEV_ACCOUNT_ID variable (which should contain your AWS Account ID for your development environment). To do this:
 - Go to the **Settings** of your cloned or forked repository in GitHub.
 - Navigate to **Secrets and Variables** > **Actions**.
 - Add the AWS_DEV_ACCOUNT_ID and any other required variables or secrets.
- c. Navigate to the Actions Page:
 - From the cloned or forked repository on GitHub, click on the **Actions** tab.
- d. Select and Run the Deploy Workflow:
 - Find the **Deploy workflow** link on the left side of the Actions page and click on it.




- On the **Deploy workflow** page, click the **Run workflow** button.
- A workflow dialog box will appear. Fill in the necessary details and click **Run workflow** to initiate the Terraform deployment process.



By following these steps, you will run the Terraform configuration directly from GitHub, leveraging GitHub Actions for automation and deployment.

3.8 Visualizing the Terraform Configuration

Below is the Terraform visualization of the Terraform configuration. It shows the resources and their dependencies, making the infrastructure setup easier to understand.



To fully view the image, open it in another tab on your browser to zoom in.

When you update the Terraform Configuration, to update the Terraform visualization, use the terraform graph command with Graphviz to generate a visual representation of the resources and their dependencies. To do this, run the following command:

terraform graph | dot -Tpng > .blog/images/terraform-visualization.png

4.0 Hands-On Kickoff: Practical Examples for Rapid Learning

4.1 Flink Applications Powered by Java on a locally running Apache Flink Cluster in Docker

Let's go!

4.2 Flink Applications Powered by Python on a locally running Apache Flink Cluster in Docker

Let's go!

4.3 Flink Applications Powered by Python on Confluent Cloud for Apache Flink

Let's go!

5.0 Powering Analytics with Apache Iceberg tables in Snowflake

After running the Flink applications, you can now leverage the data stored in the Apache Iceberg tables in Snowflake for analytics. Click here to learn more about using **Non-Managed Apache Iceberg tables** in Snowflake. This blog post will help you understand the trade-offs, operational characteristics, and integration patterns of using Apache Iceberg tables in Snowflake. These terraform configurations:

- setup-aws-s3 creates the necessary Snowflake IAM roles and policies in AWS to get started.
- setup-snowflake-service_user grants the necessary privileges to the Snowflake resources to get started.
- setup-snowflake-objects creates the necessary resources in Snowflake to get started.

6.0 Resources

Apache Flink's Core is Dataflow Programming

What is Apache Flink? — Architecture

Apache Flink Use Cases

Building Apache Flink Applications in Java

J3's techStack Lexicon

Unlocking Schema Registry Access: Granting Confluent Service Account Permissions with Terraform